

Per-instance regret difference (GRPO minus RLOO) vs sigma\_r: predicted (exact\_var) vs realized

- predicted, exact\_var ( $\theta_{\text{GRPO}}^* - \theta_{\text{true}}^*$ )
- realized (trained GRPO - RLOO, 5 seeds, 95% CI)

